

03

Programming Gen AI Applications

In this module, you learn to ...

01

Program with the PaLM REST API

02

Program Jupyter Notebooks that use the PaLM API

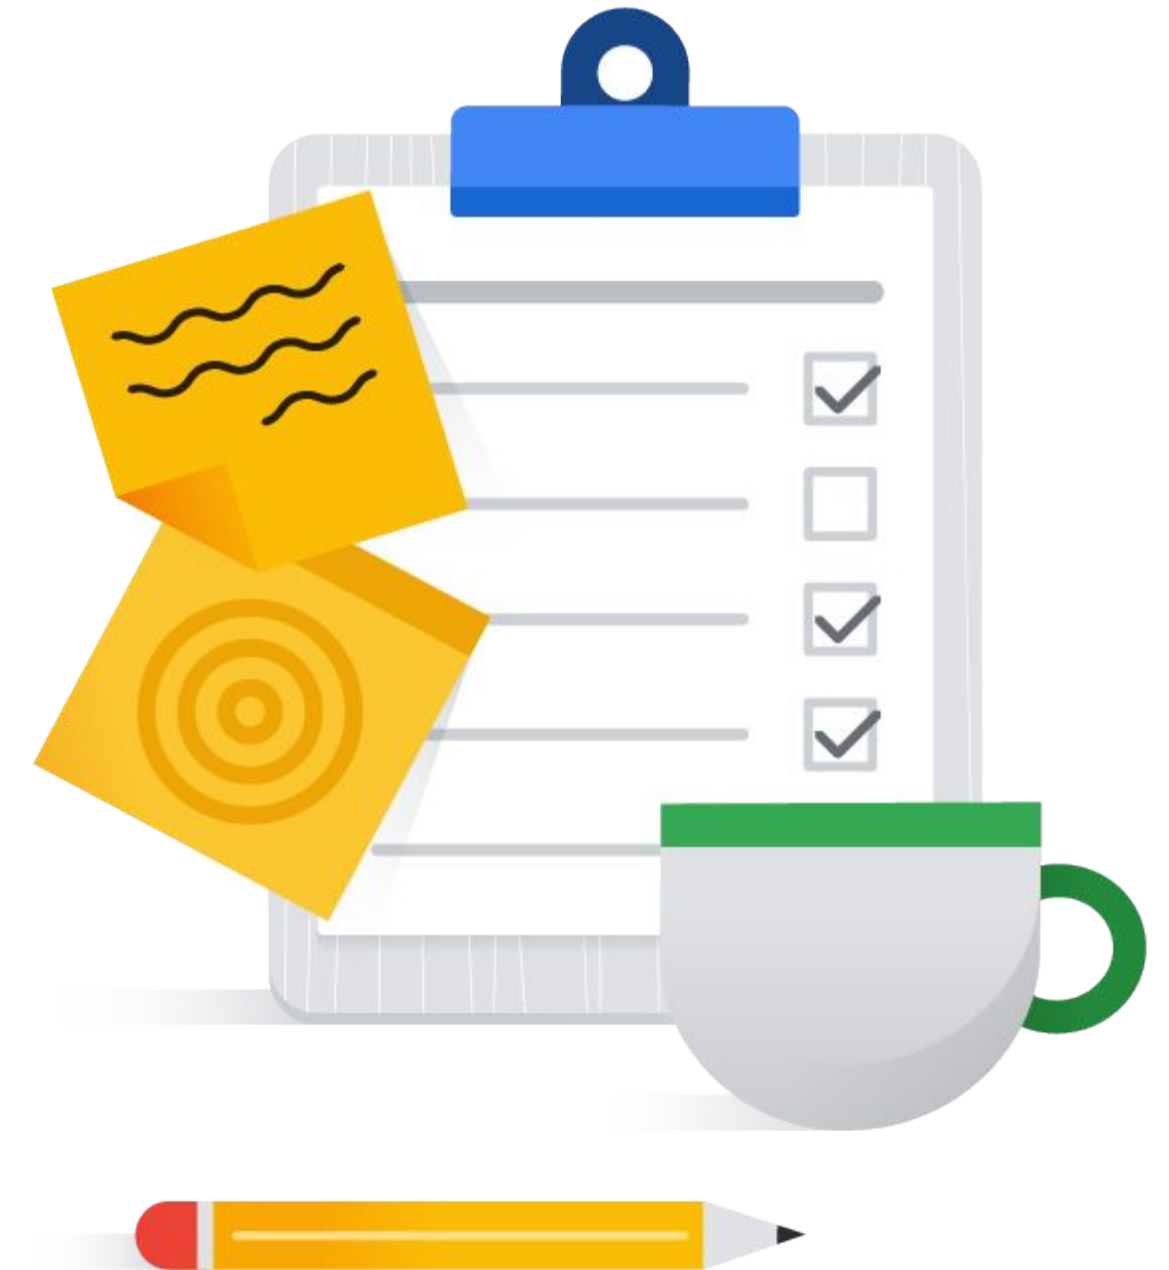
03

Add GenAI capabilities to your Python applications



Topics

- | | |
|----|--|
| 01 | Getting Started |
| 02 | Python Language API |
| 03 | Programming Text Generation Applications |
| 04 | Programming Chat Applications |
| 05 | Large Document Processing |
| 06 | Imagen and Gemini |



PaLM is a Large Language Model (LLM)

- LLMs are very sophisticated autocomplete applications
 - They learn patterns from large amounts of text
 - Use those patterns to generate text
- When generating text they calculate the next most likely tokens (words)
 - They aren't smart; it's math and statistics
- PaLM can generate text with two basic services
 - Text service for single request interactions
 - Chat service is for interactive, multi-turn interactions

To use the PaLM API, authenticate your application

- Obtain an authorization token
- Run the application using a service account

An authorization token identifies the caller of an API

- Created using the Google Cloud CLI
 - The gcloud CLI must be initialized with either a user or service account
- Set the Authorization header variable with the token generated using gcloud

```
curl \
-X POST \
-H "Authorization: Bearer $(gcloud auth print-access-token)" \
-H "Content-Type: application/json" \
"https://us-central1-aiplatform.googleapis.com/v1/projects/${PROJECT_ID}/locations/${LOCATION_ID}/publishers/google/models/${MODEL_ID}:predict" -d \
${'
    "instances": [
<<code omitted>>
```

Be careful when using the PaLM API

- Google makes available two APIs for developing with PaLM
 - One API is made available to the general public
 - For Enterprise applications, make sure to use the Google Cloud Vertex AI API
- Examine the endpoints
 - The public API uses `generativelanguage.googleapis.com`
 - The enterprise API uses `aiplatform.googleapis.com`

```
!curl https://generativelanguage.googleapis.com/v1beta2/models/text-embedding-001 \
-H 'Content-Type: application/json' \
-X POST \
-d '{ \
  "prompt": { \
    "text": "What is the capital of France?" \
  } \
}'
```

```
curl \
-X POST \
-H "Authorization: Bearer $(gcloud auth print-access-token)" \
-H "Content-Type: application/json" \
https://us-central1-aiplatform.googleapis.com/v1/projects/${PROJECT_ID}/locations/us-central1/publishers/google/models/gemini-1.0-pro-001:generateContent \
-d '{ \
  "instances": [ \
    { "prompt": "Give me ten interview questions for the role of program manager." } \
  ], \
  "parameters": { \
    "maxOutputTokens": 100, \
    "temperature": 0.2 \
  } \
}'
```


If running an application in Google Cloud, assign a service account to the runtime

- Create a service account using **IAM**
 - Assign the **Vertex AI Service Agent** role
 - Use the service account to identify the runtime
- If using Cloud Run, App Engine, or Cloud Functions, the runtime will use the Compute Engine Default Service Account by default
 - This will work as it uses the Editor role
 - However, it violates principle of least privilege
- You can also download Service Account keys to authenticate programs that use the language client libraries

The screenshot displays the Google Cloud IAM console interface. It is divided into two main sections: '1 Service account details' and '2 Grant this service account access to project (optional)'. In the first section, the 'Service account name' is 'vertex-ai-sa', and the 'Service account ID' is also 'vertex-ai-sa'. The email address is 'vertex-ai-sa@vertex-ai-dar.iam.gserviceaccount.com'. The second section, which is highlighted with a blue border, is titled 'Grant this service account access to project (optional)'. It contains a text block explaining that the service account needs permissions to complete specific actions. Below this, there is a 'Role' dropdown menu set to 'Vertex AI Service Agent', which includes a description: 'Gives Vertex AI the permissions it needs to function.' To the right of the role dropdown is an 'IAM condition (optional)' section with a '+ ADD IAM CONDITION' link and a trash icon.

1 Service account details

Service account name: vertex-ai-sa

Display name for this service account

Service account ID *: vertex-ai-sa

Email address: vertex-ai-sa@vertex-ai-dar.iam.gserviceaccount.com

2 Grant this service account access to project (optional)

Grant this service account access to vertex-ai-dar so that it has permission to complete specific actions on the resources in your project. [Learn more](#)

Role: Vertex AI Service Agent


Gives Vertex AI the permissions it needs to function.

IAM condition (optional) ?

+ ADD IAM CONDITION

Generating code with Vertex AI Studio

- In Vertex AI Studio, click the **Get Code** button
 - Returns the code in Python, iPython, and cURL



CURL code example (REST API)

Get code

PYTHON

NODE.JS

JAVA

CURL

Use the command line interface (CLI) to request a model response

1. Install Google Cloud SDK [Google Cloud SDK](#) if you haven't already or open Cloud Shell and skip to step three.
2. Run the following command to authenticate using your Google account.

```
$ gcloud auth login
```

- 3. Enter the following to request a model response

```
API_ENDPOINT="us-central1-aiplatform.googleapis.com"
PROJECT_ID="vertex-ai-dar"
MODEL_ID="gemini-pro"
LOCATION_ID="us-central1"

curl \
-X POST \
-H "Authorization: Bearer $(gcloud auth print-access-token)" \
-H "Content-Type: application/json" \
"https://$API_ENDPOINT/v1beta1/projects/$PROJECT_ID/locations/$LOCATION_ID/publishers/goc
$'{
  "contents": [
    {
      "role": "user",
      "parts": [
        {
          "text": "Tell me a funny joke"
        }
      ]
    }
  ],
  "generation_config": {
```

Java code example

Get code


PYTHON

NODE.JS


JAVA

CURL

Use this script to request a model response in your application.


1. [Set up your Java Development Environment](#) 
2. Authenticate

```
gcloud config set project PROJECT_ID
gcloud auth login ACCOUNT
```


3. Add google-cloud-vertexai as your dependency

```
<!--If you are using Maven with BOM, add the following in your pom.xml-->
<dependencyManagement>
  <dependencies>
    <dependency>
      <groupId>com.google.cloud</groupId>
      <artifactId>libraries-bom</artifactId>
      <version>26.29.0</version>
      <type>pom</type>
      <scope>import</scope>
    </dependency>
  </dependencies>
</dependencyManagement>
<dependencies>
  <dependency>
    <groupId>com.google.cloud</groupId>
    <artifactId>google-cloud-vertexai</artifactId>
  </dependency>
</dependencies>

<!--If you are using Maven without BOM, add the following to your pom.xml-->
<dependencies>
```



JavaScript code example

Get code

PYTHON

NODE.JS

JAVA

CURL

Use this script to request a model response in your application.

1. Install the Vertex AI SDK.

```
npm install https://github.com/googleapis/nodejs-vertexai
gcloud auth application-default login
```



2. Create an index.js file and add the following code:

```
const {VertexAI} = require('@google-cloud/vertexai');

// Initialize Vertex with your Cloud project and location
const vertex_ai = new VertexAI({project: 'vertex-ai-dar', location: 'us-central1'});
const model = 'gemini-pro';

// Instantiate the models
const generativeModel = vertex_ai.preview.getGenerativeModel({
  model: model,
  generation_config: {
    "max_output_tokens": 2048,
    "temperature": 0.9,
    "top_p": 1
  },
});


async function generateContent() {
  const req = {
    contents: [{role: 'user', parts: [{text: 'Tell me a funny joke'}]}],
  };

  const streamingResp = await generativeModel.generateContentStream(req);
```



Python code example


Get code PYTHON NODE.JS JAVA CURL

 OPEN NOTEBOOK

Use this script to request a model response in your application.

1. Install the Vertex AI SDK: Open a terminal window and enter the command below. You can also [install it in a virtualenv](#) .

```
!pip install --upgrade google-cloud-aiplatform
```



2. Use the following code in your application to request a model response

```
import vertexai
from vertexai.preview.generative_models import GenerativeModel, Part

def generate():
    model = GenerativeModel("gemini-pro")
    responses = model.generate_content(
        """Tell me a funny joke""",
        generation_config={
            "max_output_tokens": 2048,
            "temperature": 0.9,
            "top_p": 1
        },
        stream=True,
    )

    for response in responses:
        print(response.candidates[0].content.parts[0].text)

generate()
```



Google Cloud

Click the Open Notebook button to run the code in a Jupyter Notebook


Get code

PYTHON


NODE.JS

JAVA

CUR

 OPEN NOTEBOOK

Use this script to request a model response in your application.

1. Install the Vertex AI SDK: Open a terminal window and enter the command below. You can also [install it in a virtualenv](#)  .

```
pip install --upgrade google-cloud-aiplatform
```


REST API response

```
{ "predictions": [  
  {  
    "citationMetadata": {  
      "citations": []  
    },  
    "content": "```\nI am a programmer, \...\n```",  
    "safetyAttributes": {  
      "categories": [],  
      "blocked": false,  
      "scores": []  
    }  
  },  
  "metadata": {  
    "tokenMetadata": {  
      "inputTokenCount": {  
        "totalTokens": 8,  
        "totalBillableCharacters": 33},  
      "outputTokenCount": {  
        "totalTokens": 130,  
        "totalBillableCharacters": 355}  
    }  
  }  
}]
```

Where the response was derived

The response

Input and output tokens determine the cost

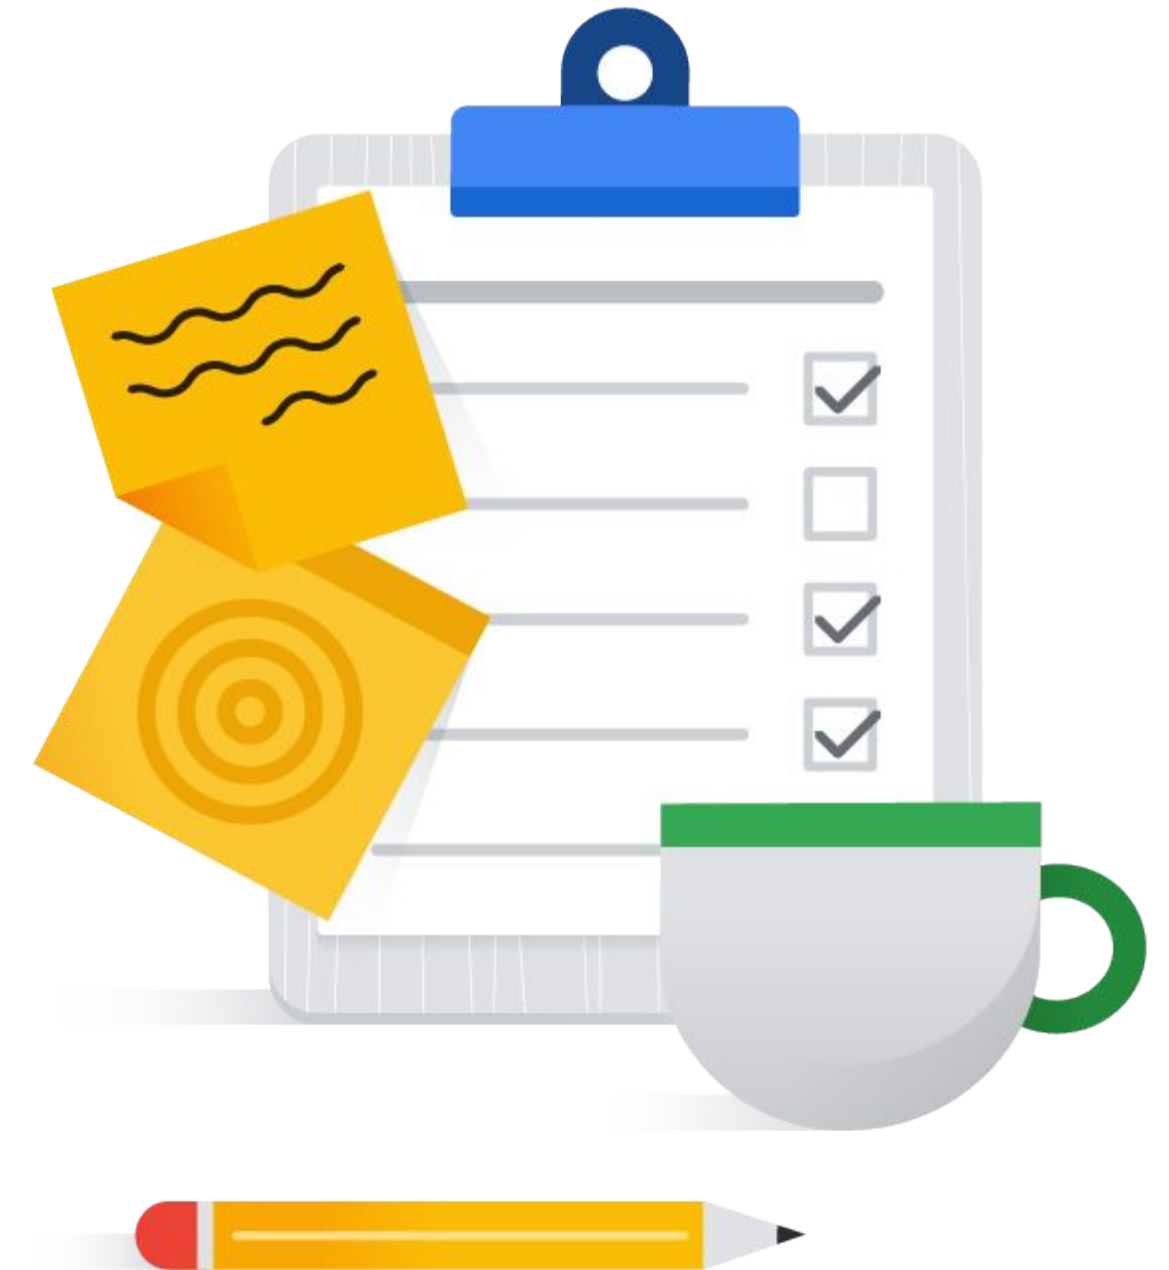
Safety attributes

- Returns an array of categories and an array of scores
 - A category is only returned if its score is greater than 0
 - Score is a value between 0.0 and 1.0
 - The higher the score the more likely the content violates that category
- You should check those values before responding to a user
 - Set a threshold where responses should be blocked

```
"safetyAttributes": {  
  "categories": [  
    "Death, Harm & Tragedy",  
    "Public Safety",  
    "Religion & Belief",  
    "War & Conflict"  
  ],  
  "blocked": false,  
  "scores": [  
    0.9,  
    0.1,  
    0.2,  
    0.8  
  ]  
}
```

Topics

01	Getting Started
02	Python Language API
03	Programming Text Generation Applications
04	Programming Chat Applications
05	Large Document Processing
06	Imagen and Gemini



Vertex AI requirements for Python

- Use pip to install Google Cloud AI Platform
 - Or add to your requirements.txt file

```
pip install google-cloud-aiplatform --upgrade
```

- Import Vertex AI and the classes required for your application

```
import vertexai  
from vertexai.language_models import TextGenerationModel
```

Basic Python code for Text Generation

```
import vertexai
from vertexai.language_models import TextGenerationModel
vertexai.init(project="your-proj-id", location="us-central1")
parameters = {
    "candidate_count": 1,
    "max_output_tokens": 1024,
    "temperature": 0.2,
    "top_p": 0.8,
    "top_k": 40
}
model = TextGenerationModel.from_pretrained("text-bison")
response = model.predict(
    """Tell me about Grace Hopper""", **parameters
)
print(f"Response from Model: {response.text}")
```

Set parameters appropriate
to your use case

Create an instance of the
model and use predict to make
the request

Get the output using the text
property of the response

Basic Python code for Chat

```
import vertexai
from vertexai.language_models import ChatModel, InputOutputTextPair

vertexai.init(project="your-proj-id", location="us-central1")
chat_model = ChatModel.from_pretrained("chat-bison")
parameters = {
    "candidate_count": 1,
    "max_output_tokens": 1024,
    "temperature": 0.2,
    "top_p": 0.8,
    "top_k": 40
}
chat = chat_model.start_chat()
response = chat.send_message("""Who was Steve Jobs?""", **parameters)
print(f"Response from Model: {response.text}")
```

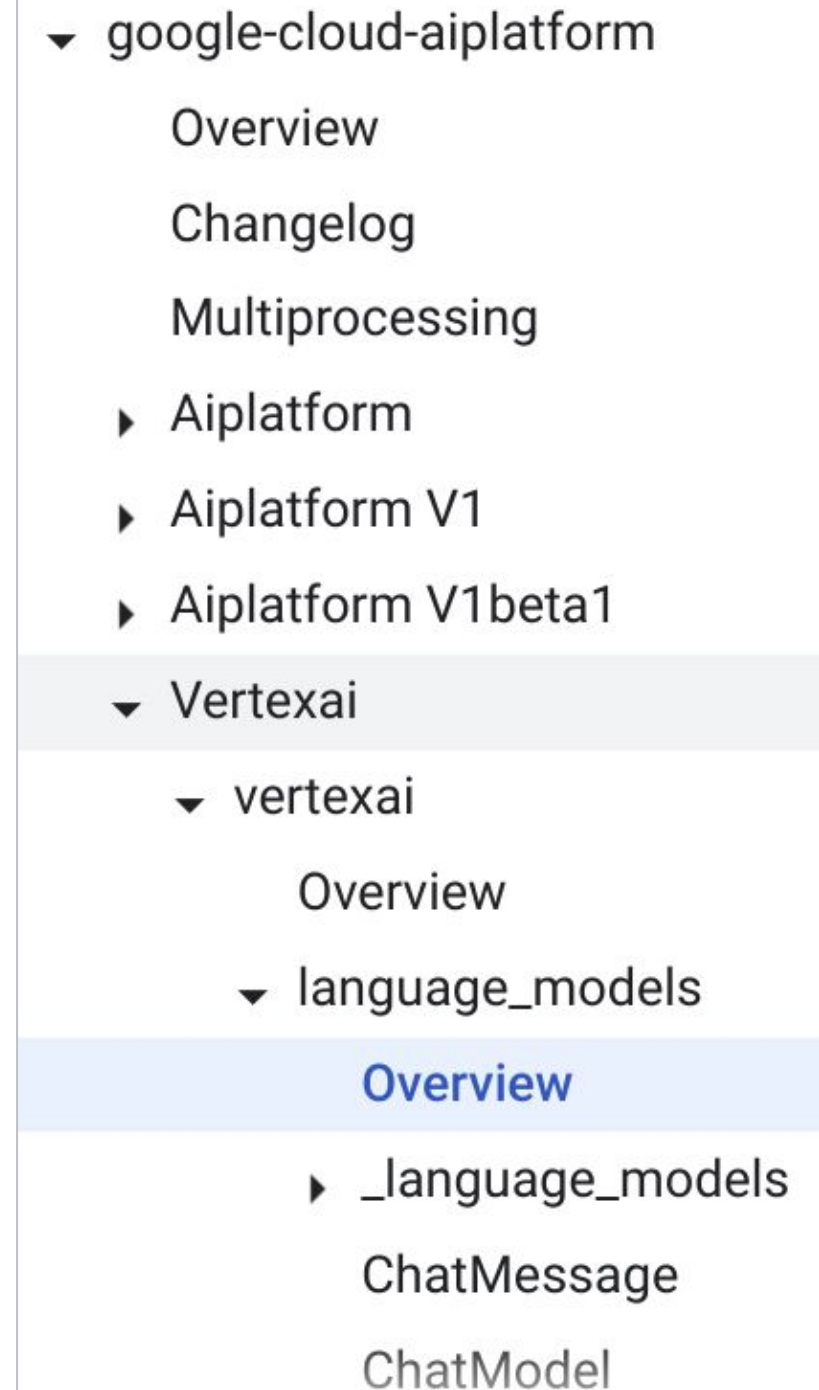
When you start a chat the history of the conversation is maintained

Request parameters

Property	Description
temperature	<ul style="list-style-type: none">● Value between 0 and 1● Controls the degree of randomness in the output● 0 is deterministic (it always returns the highest probable token)
maxOutputTokens	<ul style="list-style-type: none">● Maximum size of the response● Value between 1 and 2048
topK	<ul style="list-style-type: none">● Value between 1 and 40● Determines the number of tokens that can be chosen
topP	<ul style="list-style-type: none">● Value between 0 and 1● Tokens are selected from most probable to least until the sum of their probabilities equals the top-p value.
candidateCount	<ul style="list-style-type: none">● The number of candidate responses to return

Python API Documentation

- Drill down to:
google-cloud-aiplatform > Vertexai > vertexai > language_models
- For text generation models the important classes are:
 - TextGenerationModel
 - TextGenerationResponse
- For chat uses cases the important classes are;
 - ChatModel
 - ChatSession
 - ChatMessage
- Examples are added using the class
 - InputOutputTextPair



▼ google-cloud-aiplatform

- Overview
- Changelog
- Multiprocessing
- ▶ Aiplatform
- ▶ Aiplatform V1
- ▶ Aiplatform V1beta1
- ▼ Vertexai
 - ▼ vertexai
 - Overview
 - ▼ language_models
 - Overview**
 - ▶ _language_models
 - ChatMessage
 - ChatModel

- <https://cloud.google.com/python/docs/reference/aiplatform/latest/vertexai>

TextGenerationModel Methods

from_pretrained

Factory method to create the model instance

```
from_pretrained(model_name: str) -> vertexai._model_garden._model_garden_models.T
```


predict

```
predict(  
    prompt: str,  
    *,  
    max_output_tokens: typing.Optional[int] = 128,  
    temperature: typing.Optional[float] = None,  
    top_k: typing.Optional[int] = None,  
    top_p: typing.Optional[float] = None,  
    stop_sequences: typing.Optional[typing.List[str]] = None,  
    candidate_count: typing.Optional[int] = None  
) -> vertexai.language_models.MultiCandidateTextGenerationResponse
```

Pass in the prompt and parameters to get a model response

TextGenerationResponse contains the response from the LLM including text and safety attributes

```
TextGenerationResponse(text: str, _prediction_response: typing.Any, is_blocked: bool = False,  
safety_attributes: typing.Dict[str, float] = <factory>)
```



ChatModel Methods

from_pretrained

```
from_pretrained(model_name: str) -> vertexai._model_garden._model_garden_models.T
```

start_chat

```
start_chat(  
    *,  
    context: typing.Optional[str] = None,  
    examples: typing.Optional[  
        typing.List[vertexai.language_models.InputOutputTextPair]  
    ] = None,  
    max_output_tokens: typing.Optional[int] = None,  
    temperature: typing.Optional[float] = None,  
    top_k: typing.Optional[int] = None,  
    top_p: typing.Optional[float] = None,  
    message_history: typing.Optional[  
        typing.List[vertexai.language_models.ChatMessage]  
    ] = None,  
    stop_sequences: typing.Optional[typing.List[str]] = None  
) -> vertexai.language_models.ChatSession
```

Factory method to create the model instance



When you call the `start_chat` function a `ChatSession` is created

```
ChatSession(  
    model: vertexai.language_models.ChatModel,  
    context: typing.Optional[str] = None,  
    examples: typing.Optional[  
        typing.List[vertexai.language_models.InputOutputTextPair]  
    ] = None,  
    max_output_tokens: typing.Optional[int] = None,  
    temperature: typing.Optional[float] = None,  
    top_k: typing.Optional[int] = None,  
    top_p: typing.Optional[float] = None,  
    message_history: typing.Optional[  
        typing.List[vertexai.language_models.ChatMessage]  
    ] = None,  
    stop_sequences: typing.Optional[typing.List[str]] = None,  
)
```

Context, examples, and parameters are maintained within the `ChatSession`

The message history has to be maintained for each user

ChatSession Methods

send_message

```
send_message(  
    message: str,  
    *,  
    max_output_tokens: typing.Optional[int] = None,  
    temperature: typing.Optional[float] = None,  
    top_k: typing.Optional[int] = None,  
    top_p: typing.Optional[float] = None,  
    stop_sequences: typing.Optional[typing.List[str]] = None,  
    candidate_count: typing.Optional[int] = None,  
    grounding_source: typing.Optional[  
        typing.Union[  
            vertexai.language_models._language_models.WebSearch,  
            vertexai.language_models._language_models.VertexAISearch,  
        ]  
    ] = None  
) -> vertexai.language_models.MultiCandidateTextGenerationResponse
```

The ChatSession message_history property is a collection of ChatMessage objects

```
ChatMessage(content: str, author: str)
```

Examples are added using InputOutputTextPair objects

```
InputOutputTextPair(input_text: str, output_text: str)
```

Streaming responses can make the user interface more responsive

- Streaming with text generation use the `predict_streaming` function
 - Returns a collection of responses

```
responses = model.predict_streaming("Tell me about Steve Jobs", **parameters)

for response in responses:
    print(response.text)
```

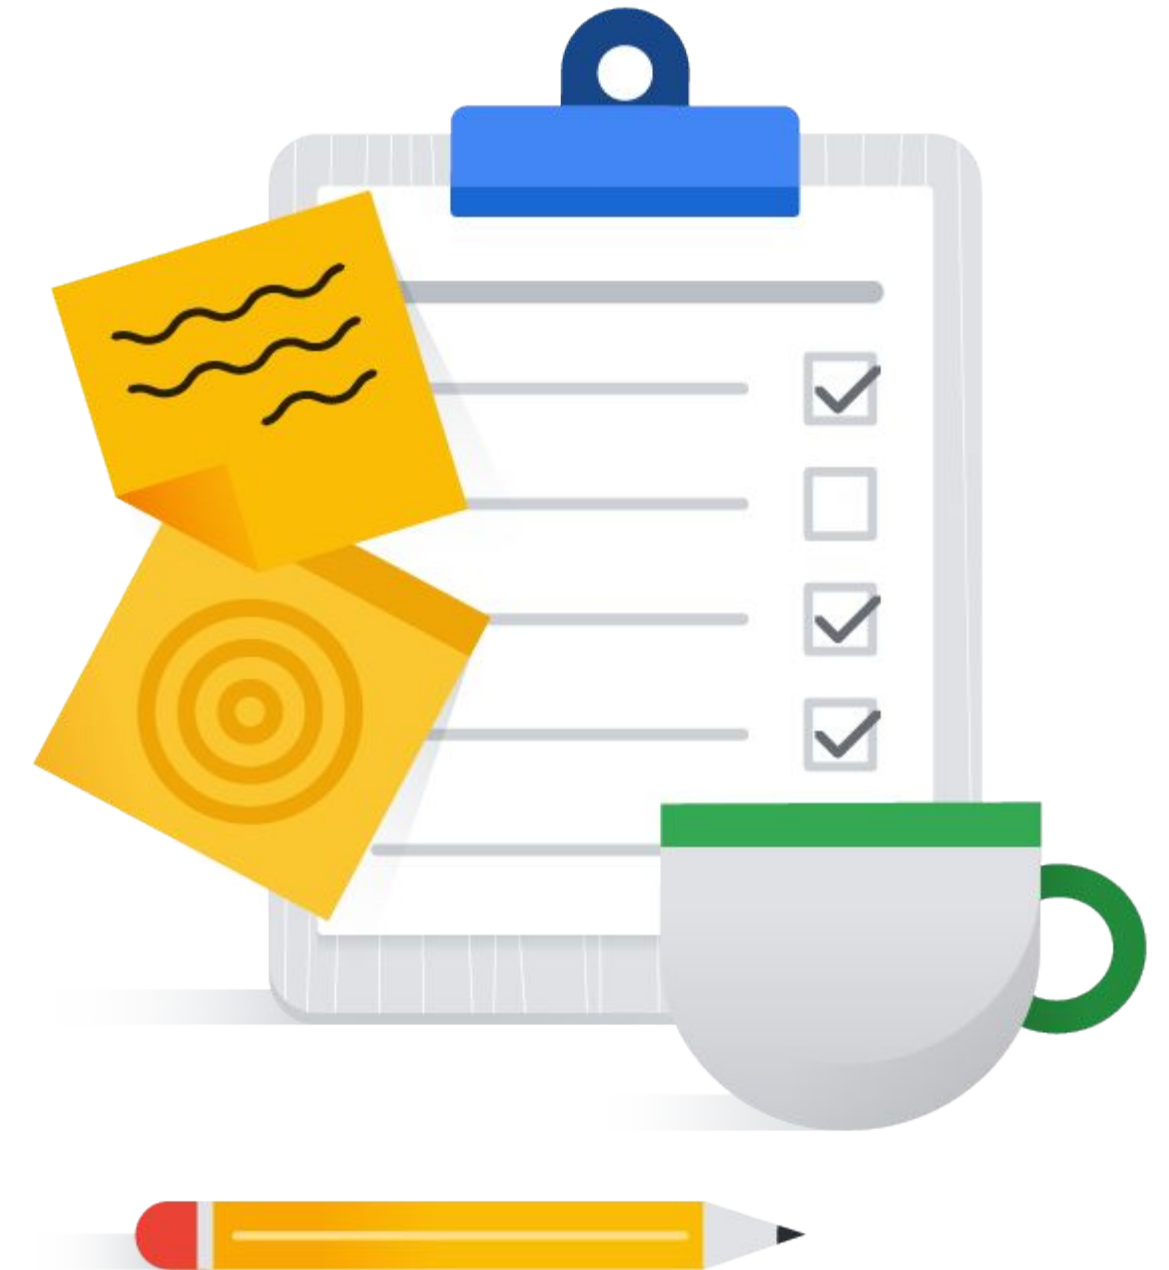
- Streaming with a chat session use the `send_message_streaming` function

```
responses = chat.send_message_streaming("Tell me about Grace Hopper", **parameters)

for response in responses:
    print(response.text)
```

Topics

01	Getting Started
02	Python Language API
03	Programming Text Generation Applications
04	Programming Chat Applications
05	Large Document Processing
06	Imagen and Gemini



Python Flask Website example

- This is an example of using the text service with the PaLM API
 - Even though you may ask many questions, each one is independent
- Context must be added to tell the PaLM API to emulate a barista
- The coding is simple as you are just submitting an HTML form and making a request to the PaLM API for a response
 - The response is displayed on the screen

CoffeeBot

CoffeeBot
Your friendly online BaristaAI

I am CoffeeBot, a barista and expert on all things related to coffee and tea. I can help you find the perfect coffee or tea for your taste, and I can also teach you how to make your own coffee and tea drinks at home.

Ask CoffeeBot:

Submit

CoffeeBot

CoffeeBot
Your friendly online BaristaAI

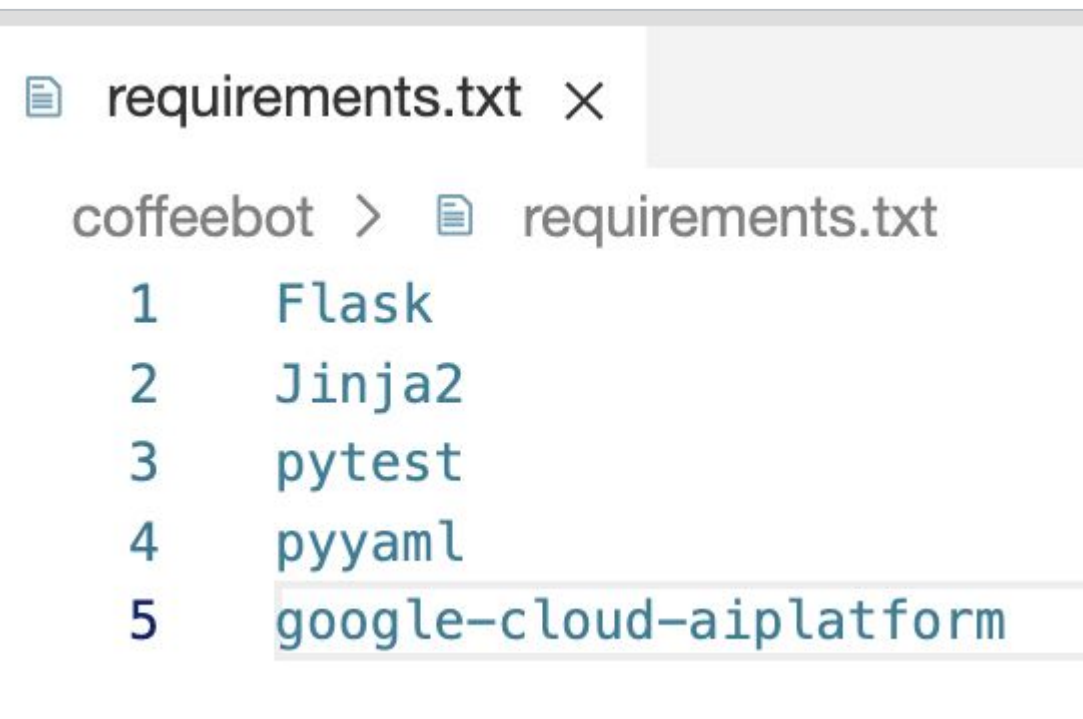
To make a latte, you will need: * 2 shots of espresso * 6 ounces of steamed milk * 1 tablespoon of foamed milk 1. Brew the espresso shots. 2. Steam the milk until it is hot and frothy. 3. Pour the espresso into a latte glass. 4. Add the steamed milk to the espresso. 5. Top with the foamed milk. 6. Enjoy!

Ask CoffeeBot:

Submit

Add the Python requirements

- Add Google Cloud AI Platform to the requirements.txt file
- Add the required imports at the top of the code file



The screenshot shows a code editor window with a tab labeled 'requirements.txt'. The file path 'coffeobot > requirements.txt' is visible. The file contains five lines of requirements, with the fifth line, 'google-cloud-aiplatform', highlighted.

```
requirements.txt
1 Flask
2 Jinja2
3 pytest
4 pyyaml
5 google-cloud-aiplatform
```

```
from flask import Flask, render_template, request
import os
import vertexai
from vertexai.language_models import TextGenerationModel
```


Handling web requests in Flask

- The default route will handle HTTP posts and gets
 - Post means a question was submitted from the HTML form
 - Get means there is no question (have CoffeeBot introduce itself)
- The code for using the PaLM API is in the `get_response()` function

```
@app.route("/", methods = ['POST', 'GET'])
def main():
    if request.method == 'POST':
        input = request.form['input']
        response = get_response(input)
    else:
        input = ""
        response = get_response("Who are you and what can you do?")

    model = {"title": "CoffeeBot", "message": response, "input": input}
    return render_template('index.html', model=model)
```

Making a request to the PaLM API

```
def get_response(input):  
    vertexai.init(project="vertext-ai-dar", location="us-central1")  
    parameters = {  
        "temperature": 0.8,  
        "max_output_tokens": 256,  
        "top_p": 0.8,  
        "top_k": 40  
    }  
  
    model = TextGenerationModel.from_pretrained("text-bison@001")
```

Initialize the API
and set up the
parameters

Create the
model using the
correct version
of the PaLM API

<< CODE CONTINUED ON NEXT SLIDE >>

Making a request to the PaLM API (continued)

```
def get_response(input):
```

```
<< CODE CONTINUED FROM PREVIOUS SLIDE >>
```

```
model = TextGenerationModel.from_pretrained("text-bison@001")  
request = """Your name is CoffeeBot. You are a barista and expert on  
all things related to coffee and tea..
```

The context tells
the API to
emulate a barista

```
input: {}
```

The input is the
question

```
output:
```

```
"""
```

```
response = model.predict(  
    request.format(input),  
    **parameters
```

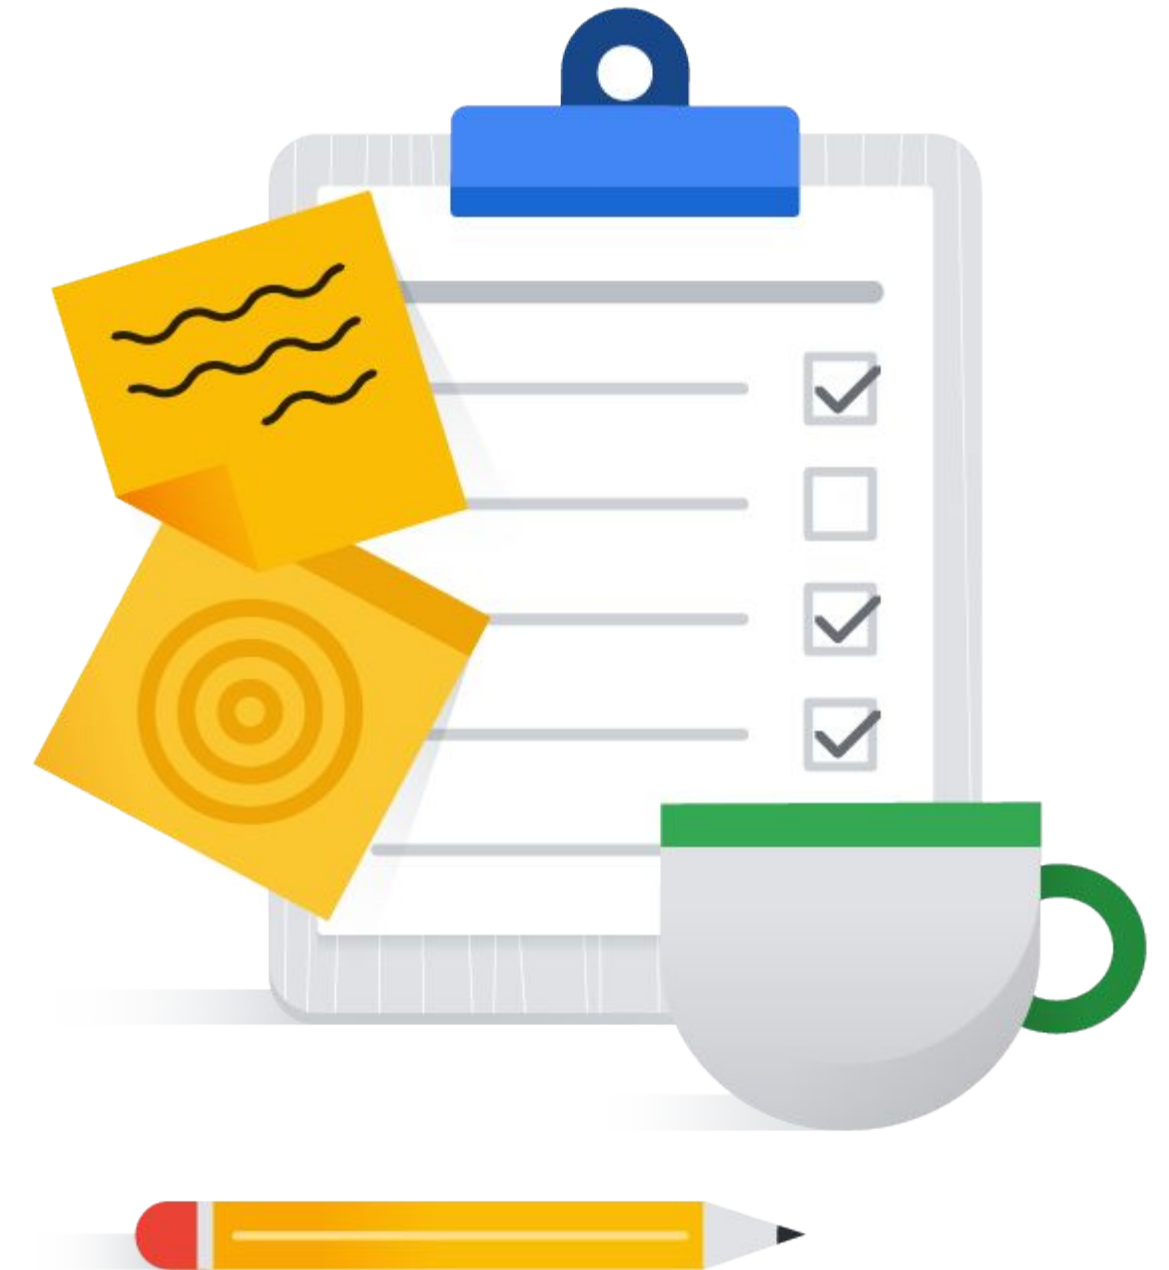
Call the `model.predict()`
function to send the request

```
)  
return response
```

The format function injects the
question into the prompt

Topics

- | | |
|----|--|
| 01 | Getting Started |
| 02 | Python Language API |
| 03 | Programming Text Generation Applications |
| 04 | Programming Chat Applications |
| 05 | Large Document Processing |
| 06 | Imagen and Gemini |



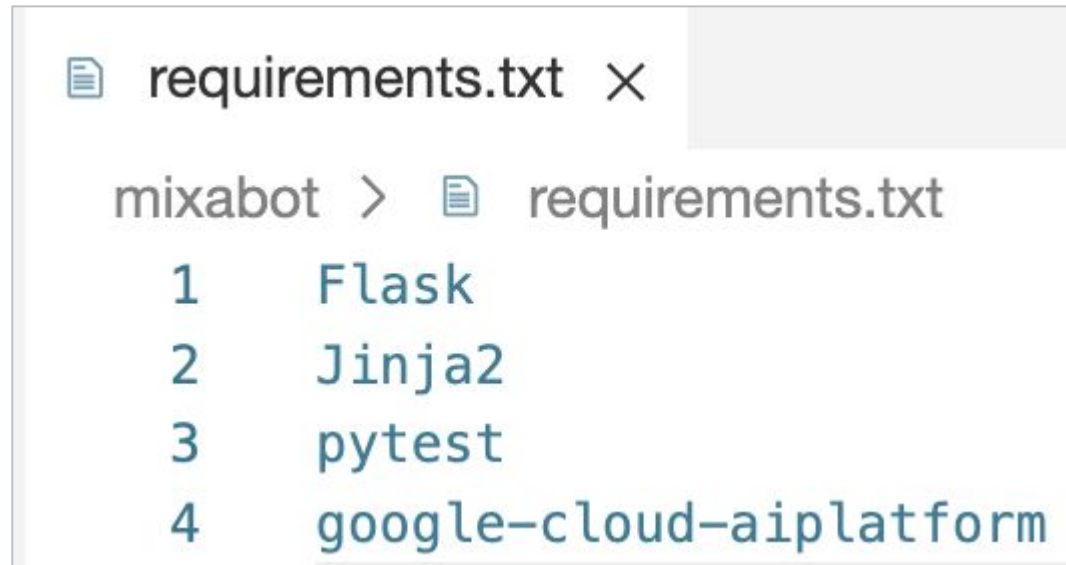
Python Flask Website example

- This is an example of using the chat service with the PaLM API
 - The chat will remember the history of the conversation
- Context must be added to tell the PaLM API to emulate a customer service agent for the service station
- The coding is simple as you are just submitting an HTML form and making a request to the PaLM API for a response
 - The response is displayed on the screen

The image shows two overlapping screenshots of a web chat interface. The top screenshot shows the initial state of the chat. It has a header 'ChET' in blue. Below it is a large grey box with the text 'ChET' in bold and 'Your friendly online mechanic from North Point Gas and All' below it. At the bottom, there is a green box with the text 'I am ChET, the Chatty and All. I can help you'. Below that is a form with the label 'Ask ChET:', an input field, and a blue 'Submit' button. The bottom screenshot shows the chat after a user has submitted a message. The green box now contains the text 'Have you tried turning the car off and back on again?'. The input field now contains the text 'My car makes a noise when I start the engine'. The 'Submit' button is still present.

Add the Python requirements

- Add Google Cloud AI Platform to the requirements.txt file
- Add the required imports at the top of the code file

A screenshot of a code editor window. The title bar shows a document icon, the filename 'requirements.txt', and a close button 'X'. The editor content shows a directory path 'mixabot >' followed by a document icon and the filename 'requirements.txt'. Below this, there is a list of four items, each with a line number on the left: '1 Flask', '2 Jinja2', '3 pytest', and '4 google-cloud-aiplatform'. The text is in a monospaced font, with the package names in a blue color.

```
requirements.txt X  
mixabot > requirements.txt  
1 Flask  
2 Jinja2  
3 pytest  
4 google-cloud-aiplatform
```

```
from flask import Flask, render_template, request  
import os  
import vertexai  
From vertexai.preview.language_models import ChatModel, InputOutputTextPair,  
ChatMessage
```


Initializing the Chat session

```
vertexai.init(location="us-central1")
chat_model = ChatModel.from_pretrained("chat-bison@001")
parameters = {
    "temperature": TEMPERATURE,
    "max_output_tokens": MAX_OUTPUT_TOKENS,
    "top_p": TOP_P,
    "top_k": TOP_K
}
examples=[
    InputOutputTextPair(
        input_text="""When I turn my car on, there is a clicking noise. """,
        output_text="""Did you try turning the engine off and back on again?""")
]
chat = chat_model.start_chat(context=CONTEXT, examples=examples, **parameters)
```

Initialize the API and set up the parameters

Add examples

Start the chat

Managing User Sessions

- In a chat, the history of the conversation needs to be maintained per user
 - The ChatSession object has a message_history property
- Create a session variable with the history for each user
 - Reinitialize the chat with every request setting the message history property
- In Python Flask, sessions are stored in the client browser, so this is a scalable solution

```
response = chat.send_message(input)
session["chat_history"] = chat.message_history
```

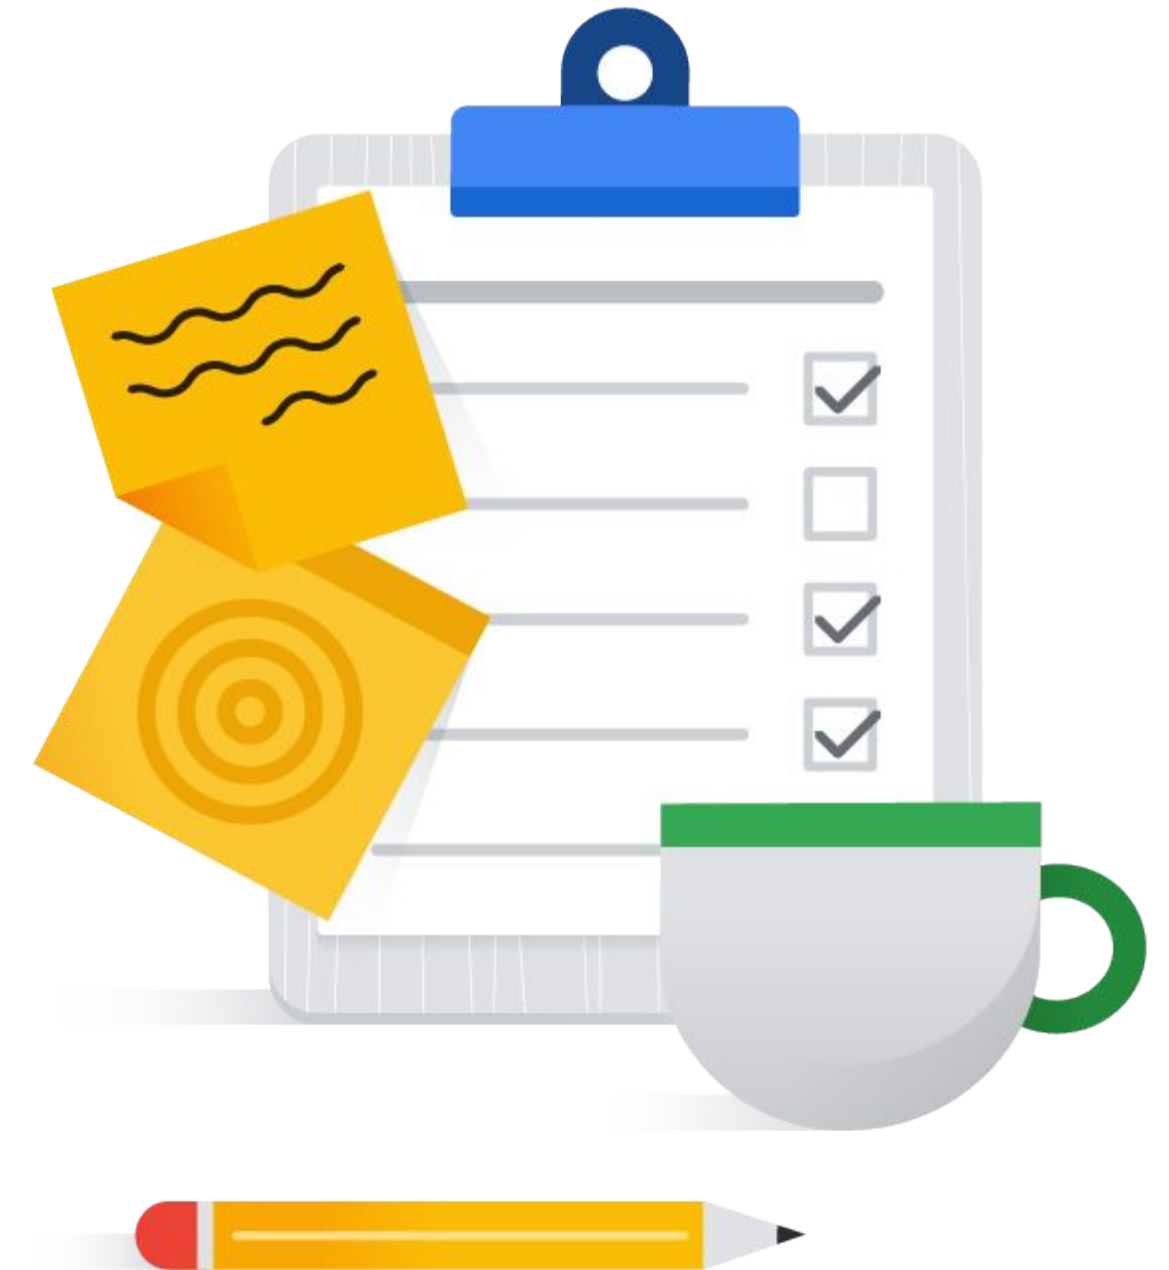
Need to convert the history
in a session variable into
ChatMessage objects



```
if 'chat_history' in session:
    chat_history = [ChatMessage(content=items["content"], author=items["author"]) for
                    items in session["chat_history"]]
    parameters["message_history"] = chat_history
```

Topics

01	Getting Started
02	Python Language API
03	Programming Text Generation Applications
04	Programming Chat Applications
05	Large Document Processing
06	Imagen and Gemini



What are tokens?

- When text is sent to an LLM it is split into tokens
 - A token represents an idea that the large language model understands
- Tokens are on average about 4-5 characters long
 - A short word may be converted into a single token
 - Large words might use multiple tokens
 - Punctuation is represented as tokens
- When they are being processed, tokens are converted into numeric arrays called embeddings
- The model generates output embeddings which are converted back into tokens and returned to the caller
- A token is approximately four characters (100 tokens correspond to roughly 60-80 words)

There is a limit to the number of input and output tokens that can be processed by the model

- Limits change based on model version and will likely increase over time
 - The `text-bison-32k` and `chat-bison-32k` support 32,000 tokens in the request-response
- If you want the model to summarize documents that exceed the token limit you need to split the operation into multiple calls
 - Like a map-reduce operation
 - Divide the document into pieces
 - Summarize each piece
 - Summarize the summaries
- With chat uses cases the entire conversation is sent with each request
 - Be careful that the conversation doesn't exceed the limit
 - You may need to trim the history at some point
 - Alternatively, you can have the LLM summarize the conversation thus far, then send only the summary as context

Summarizing documents is a common task for large language models

- This seems simple enough:
 - Retrieve the document to be summarized
 - Add the document to the prompt
 - Ask the LLM to summarize it
- Problems:
 - What if you have many small documents to summarize?
 - What if the data is too large for a single request?
 - What happens if you exceed the quota of requests?

Retrieving external data

- In the example below, the Python `urllib` package is used to download a pdf file and store it locally in a folder “data”

```
data_folder = "data"
Path(data_folder).mkdir(parents=True, exist_ok=True)

pdf_url =
"https://services.google.com/fh/files/misc/practitioners_guide_to_mlops_whitepaper.pdf"
pdf_file = Path(data_folder, pdf_url.split("/")[-1])

urllib.request.urlretrieve(pdf_url, pdf_file)
```

Processing the document

- Here, the document is divided into pages

```
reader = PyPDF2.PdfReader(pdf_file)
pages = reader.pages

# Print three pages from the pdf
for i in range(2):
    text = pages[i].extract_text().strip()
    print("_____")
    print(f"Page {i}: {text} \n\n")
```

```
Page 0: Practitioners guide to MLOps:
A framework for continuous
delivery and automation of
machine learning.White paper
May 2021
Authors:
Khalid Salama,
Jarek Kazmierczak,
Donna Schut
```

Stuffing means you are combining content from multiple documents or pages

- By combining documents you can reduce the number of calls to the model
 - You also can get a single summary of more than 1 thing
- In the example below, the text from all the pages is concatenated

```
reader = PyPDF2.PdfReader(pdf_file)
pages = reader.pages
concatenated_text = ""

for page in tqdm(pages):
    text = page.extract_text().strip()
    concatenated_text += text
```

In the example below, the first 19,000 words are passed to the model for summarization

- In this case, the entire document cannot be processed in one request

```
prompt_template = """
    Write a concise summary of the following text delimited by triple backquotes.
    Return your response in bullet points which covers the key points of the text.
    ```{text}```
 BULLET POINT SUMMARY:
 """

prompt = prompt_template.format(text=concatenated_text[:19000])

Use the model to summarize the text using the prompt
summary = generation_model.predict(prompt=prompt, max_output_tokens=1024).text
```

# The MapReduce algorithm can be used with large docs

- Steps:
  - Divide the document into chunks
  - Summarize the chunks (Map)
  - Combine the summaries
  - Summarize the summaries (Reduce)

# The Map step summarizes each page and adds each summary to a collection

```
reader = PyPDF2.PdfReader(pdf_file)
pages = reader.pages

initial_summary = []
for page in tqdm(pages):
 text = page.extract_text().strip()
 prompt = initial_prompt_template.format(text=text)
 summary = model_with_limit_and_backoff(prompt=prompt, max_output_tokens=1024).text
 initial_summary.append(summary)
```



# In the Reduce step, combine all the summaries and summarize those

- Be careful that the combined summaries don't exceed the the maximum length of a request

```
def reduce(initial_summary, prompt_template):
 # Concatenate the summaries from the initial step
 concat_summary = "\n".join(initial_summary)

 prompt = prompt_template.format(text=concat_summary)
 summary = model_with_limit_and_backoff(prompt=prompt, max_output_tokens=1024).text
 return summary

summary = reduce(initial_summary, final_prompt_template)
```

Notice, the call to the LLM is being made through a helper function that limits the frequency of requests

# MapReduce introduces a rate limiting problem

- Quotas vary by model, project, and region
  - Likely, quotas will change over time
- Run the following gcloud command to see your quota by model

```
gcloud alpha services quota list --service=aiplatform.googleapis.com
--consumer=projects/vertex-ai-dar
--filter=metric=aiplatform.googleapis.com/online_prediction_requests_per_base_model
```

```
effectiveLimit: 10
- defaultLimit: '60'
 dimensions:
 base_model: text-bison
 region: us-central1
 effectiveLimit: '60'
```

# Rate limiting code

```
CALL_LIMIT = 20 # Number of calls to allow within a period
ONE_MINUTE = 60 # One minute in seconds
FIVE_MINUTE = 5 * ONE_MINUTE

A function to print a message when the function is retrying
def backoff_hdlr(details):
 print(
 "Backing off {} seconds after {} tries".format(
 details["wait"], details["tries"]
)
)

Continued on next slide
```

# Rate limiting code (continued)


```
@backoff.on_exception(# Retry with exponential backoff strategy when exceptions occur
 backoff.expo,
 (
 exceptions.ResourceExhausted,
 ratelimit.RateLimitException,
), # Exceptions to retry on
 max_time=FIVE_MINUTE,
 on_backoff=backoff_hdlr, # Function to call when retrying)

@ratelimit.limits(# Limit the number of calls to the model per minute
 calls=CALL_LIMIT, period=ONE_MINUTE)

This function will call the `generation_model.predict` function, but it will retry if
defined exceptions occur.
def model_with_limit_and_backoff(**kwargs):
 return generation_model.predict(**kwargs)
```

# Topics

- |    |                                          |
|----|------------------------------------------|
| 01 | Getting Started                          |
| 02 | Python Language API                      |
| 03 | Programming Text Generation Applications |
| 04 | Programming Chat Applications            |
| 05 | Large Document Processing                |
| 06 | Imagen and Gemini                        |




# Imagen is Google's foundation model for computer vision tasks



Image generation

---



Visual Q&A  
Image Captioning

---



# Using Imagen for image generation (in preview)

```
from vertexai.preview.vision_models import ImageQnAModel, Image

model = ImageGenerationModel.from_pretrained("imagegeneration@002")
response = model.generate_images(
 prompt="Australian Shepherd herding sheep in a field, focus on the dog",
 # Optional:
 number_of_images=1
)
response[0].show()
response[0].save("shepherd.png")
```



# Using Imagen for image captioning

```
from vertexai.vision_models import ImageCaptioningModel, Image

model = ImageCaptioningModel.from_pretrained("imagetext@001")
image = Image.load_from_file("shepherd.png")
captions = model.get_captions(
 image=image,
 number_of_results=3,
 language="en",
)
for caption in captions:
 print(caption)
```

```
a dog is jumping over a sheep in a field
a dog jumping over a sheep in a field
a dog is jumping over a sheep in a grassy field
```





# Using Imagen for image Q&A

```
from vertexai.vision_models import ImageQnAModel, Image

model = ImageQnAModel.from_pretrained("imagetext@001")
image = Image.load_from_file("shepherd.png")
answers = model.ask_question(
 image=image,
 question="what kind of dog is in this picture?",
 # Optional:
 number_of_results=3,
)
print(answers)
```

```
['border collie', 'shepherd', 'collie']
```



# Using Gemini for Text Generation

```
import vertexai
from vertexai.preview.generative_models import GenerativeModel, Part

def generate():
 model = GenerativeModel("gemini-pro")
 responses = model.generate_content(
 """Do Border Collies make good pets?""",
 generation_config={
 "max_output_tokens": 2048,
 "temperature": 0.9,
 "top_p": 1}, stream=True,)

 for response in responses:
 print(response.candidates[0].content.parts[0].text)

generate()
```


The API is different when using Gemini. It is not just a matter of specifying a different model.

# Using Gemini for Chat

```
def multiturn_generate_content():
 config = {
 "max_output_tokens": 2048,
 "temperature": 0.9,
 "top_p": 1
 }
 model = GenerativeModel("gemini-pro")
 chat = model.start_chat()
 print(chat.send_message("Hi", generation_config=config))

multiturn_generate_content()
```

With Gemini, the same model is used for text generation and chat apps.



# Using Gemini for Vision

```
from vertexai.preview.generative_models import GenerativeModel, Image


multimodal_model = GenerativeModel("gemini-pro-vision")

image = Image.load_from_file("image.jpg")
prompt = "Describe this image?"

contents = [image, prompt]
responses = multimodal_model.generate_content(contents, stream=True)

for response in responses:
 print(response.text, end="")
```

The prompt along with the image(s) and/or video(s) are passed to the model






# Lab

🕒 30 min ⚙️

## Lab: Getting Started with the PaLM API for Chatbots



# In this module, you learned to ...

01

Program with the PaLM REST API

02


Program Jupyter Notebooks that use the PaLM API

03

Add GenAI capabilities to your Python applications



# Questions and answers



# Quiz question

Which of the following methods can you use to authorize PaLM API requests from an application?

- A: Obtain an authorization token and pass it in the header of the request
- B: Assign a service account to your application runtime environment
- C: Use a service account key
- D: All of the above depending on the specific use case

# Quiz question

Which of the following methods can you use to authorize PaLM API requests from an application?

- A: Obtain an authorization token and pass it in the header of the request
- B: Assign a service account to your application runtime environment
- C: Use a service account key
- D: All of the above depending on the specific use case

# Quiz question

What is the main difference between a Text Generation and Chat program?

- A: Text generation uses a large language model, chat does not
- B: Chat uses a large language model, text generation does not
- C: With text generation you have to maintain the history
- D: With chat you have to maintain the history



# Quiz question

What is the main difference between a Text Generation and Chat program?

- A: Text generation uses a large language model, chat does not
- B: Chat uses a large language model, text generation does not
- C: With text generation you have to maintain the history
- D: With chat you have to maintain the history

Google Cloud